

# ReduLink: Authenticated Reference Substitution for Redundancy-Suppressed Transfers over Encrypted QUIC Streams

Michél Nguyen

University of the People | ORCID: 0000-0001-6834-4422

## Abstract

Encrypted transports such as QUIC deliberately hide payload bytes from the network, which disables the transparent in-network redundancy elimination that wide-area optimizers have historically relied on. ReduLink is an endpoint-controlled representation layer that recovers redundancy savings for receivers that already hold related bytes, without breaking encryption: it replaces repeated payload chunks with compact references that are individually authenticated against epoch, scope, stream identifier, reconstructed offset, length, nonce, and chunk identity. The contribution is not a new chunk-matching algorithm and not a claim of faster physical links; it is authenticated, scoped, QUIC-compatible reference substitution with explicit reconstruction, privacy scoping, and fail-closed semantics.

The artifact maps compact binary ReduLink messages into TLS-protected aioquic QUIC streams, implements semantic MISS/FULL repair, derives keys with an exporter-style schedule, and tests replay and tamper rejection. We evaluate against fixed-block reuse, gzip/zstd, and real rsync over deterministic journal fixtures, hash-pinned public source releases, object-aligned public-release transfers, and a layer-like positive case, and we report native QUIC stream experiments under datagram loss, scaling, repeated trials, block-size sensitivity, component costs, competing flow and bottleneck-emulation fairness, and conservative wire-byte accounting. We add an independent, hash-pinned package version-pair study, a formal security analysis of the reference-authentication construction, a measured full-duplex shared-bottleneck path emulation, and a native miss-rate sensitivity sweep. The emulation corrects an earlier half-duplex measurement of ours: on byte-stable content ReduLink's fewer encoded bytes do shorten completion time on a constrained shared link (0.65x of the raw flow at 5 Mbps, converging as the link uncongests), but on a real repair-heavy layered workload the MISS/repair round trips can instead make it slower (up to 1.27x on a low-bandwidth high-RTT path), so the completion-time effect depends on both the link and the miss rate.

Quantitatively, on object-aligned public release transfers ReduLink reaches 1.94x to 9.70x effective stream-payload reduction under the offline model's framing (1.8x to 7.9x when re-priced at the artifact's measured 108-byte binary wire framing), and 21.37x (15.3x re-priced) on a layer-like positive case; the authenticated variant costs a further 3 to 11 percent. We also report the baselines that beat ReduLink on bytes: rsync wins by up to 73x on ordinary source-tree updates, and a zstd dictionary-delta baseline (the Compression Dictionary Transport analog) wins by 65x to 1,785x whenever the receiver retains the exact prior byte stream and a codec delta is acceptable. ReduLink's contribution is therefore not byte-optimality but authenticated, object-granular, fail-closed reference substitution over a chunk dictionary. Native aioquic experiments confirm byte-exact reconstruction with deterministic stream multipliers under zero and periodic datagram loss, and the accounting experiments confirm that congestion and fairness are charged on encoded bytes, not reconstructed bytes. The result is a clear deployment boundary: simple reuse or rsync wins for file-tree synchronization, while ReduLink is preferable for object-aligned or layer-like warm-state transfers where an authenticated, stream-compatible reference is the desired abstraction.

**Keywords:** QUIC; HTTP/3; redundancy elimination; data deduplication; content-defined chunking; shared dictionaries; authenticated references; encrypted transport; reproducible evaluation.

## 1. Introduction

Encrypted transports restrict transparent in-network redundancy elimination. QUIC combines UDP transport, TLS security, stream multiplexing, loss recovery, and congestion control, which makes middlebox rewriting of payloads the wrong abstraction for most public or end-to-end encrypted deployments [6-10]. A WAN optimizer can only suppress redundancy that it can see, and a correctly deployed QUIC connection ensures that it sees only ciphertext. The redundancy, however, has not disappeared: registries, content-delivery networks, software-update services, and backup systems still send byte-stable objects to receivers that already hold closely related bytes from a previous transfer.

The practical question is therefore not whether repeated bytes can be replaced by references. Fixed-block reuse, rsync, and low-bandwidth file systems already demonstrate that repeated bytes can be avoided [1-5], and HTTP delta and shared-dictionary mechanisms show that receiver-held bytes can be reused as a dictionary on the web [19-22]. The question this paper addresses is whether a reference-substitution mechanism can be made explicit, individually authenticated, privacy-scoped, and compatible with encrypted endpoint streams, without claiming a universal accelerator and without asking the application to become a file-tree synchronization protocol. Figure 1 shows the resulting architecture: cooperating endpoints that already control the plaintext decide, per chunk, whether to send bytes or an authenticated reference, and the receiver reconstructs or fails closed.

We use one accounting vocabulary throughout. The effective multiplier is reconstructed or input bytes divided by encoded bytes at a stated accounting layer. Stream-payload multipliers exclude UDP/IP/link overhead; UDP-estimated multipliers add a local IPv4/UDP estimate; and congestion fairness, where claimed, applies to encoded bytes, not reconstructed bytes. This separation is deliberate: ReduLink changes how many application bytes a given number of wire bytes can reconstruct, not the physical line rate.

This paper investigates three questions. RQ1: can reference substitution over encrypted QUIC streams be made explicit, authenticated, and privacy-scoped without custom transport-layer changes? RQ2: under which workload shapes does authenticated reference substitution save bytes, and where do simpler fixed-block reuse, compression, or rsync win? RQ3: what is the component-level cost of the authentication and accounting machinery, and how stable and fair are the native QUIC results across repeated runs, scales, and emulated bottlenecks? Section 7 answers RQ2, Sections 8 and 10 answer RQ3, and Sections 4 and 5 answer RQ1.

## 2. Contributions and Claim Boundary

- An authenticated ReduLink reference format and validation model for endpoint-controlled encrypted streams, with per-frame binding of epoch, scope, stream id, reconstructed offset, length, nonce, and chunk identity.
- A compact binary mapping carried inside native aioquic QUIC streams, with byte-exact semantic MISS/FULL repair, repeated trials, and a payload-scaling experiment.
- An exporter-style artifact key schedule and security tests for wrong secret, scope, epoch, stream id, offset, length, replay, and tampering, including authenticated-UDP tamper and replay probes.
- A comparison against fixed-block reuse, gzip/zstd, and real rsync, including cases where these baselines win decisively, plus an explicit contrast with HTTP delta and shared-dictionary transport.
- External public source-release negative evidence, object-aligned public source-release workloads, and a public Redis-derived layer-like positive case, all hash-pinned and reproducible.
- Block-size sensitivity, component-cost measurements, competing-flow and bottleneck-emulation fairness analysis, and conservative accounting-layer separation.

The central claim is conditional. ReduLink helps when byte-stable chunks recur across warm endpoint state and authenticated reference substitution is preferable to a file-oriented delta protocol or local compression. It is not a replacement for rsync or compression, and the artifact is a native QUIC stream mapping rather than a custom QUIC extension-frame implementation. A practical example is a registry, CDN, backup, or update service that repeatedly sends object records to the same endpoint or tenant while preserving end-to-end QUIC encryption. In that setting the server can reference receiver-known objects or chunks without exposing plaintext to a middlebox; the price is authentication and metadata overhead, and the benefit is explicit context binding, safe miss repair, and a stream-compatible representation.

## 3. Background and Related Work

### 3.1 Network redundancy elimination

Spring and Wetherall introduced protocol-independent redundancy elimination, suppressing repeated byte ranges on a link by caching recently seen content at both ends [1]. Subsequent work generalized this to network-wide and coordinated redundancy elimination (SmartRE) and to enterprise and end-system services (EndRE), and measured how much redundancy real traffic contains [2-4]. These systems are powerful but assume a vantage point that can

observe or terminate plaintext. ReduLink targets the opposite regime, where the transport is end-to-end encrypted and only the endpoints can act, so the redundancy decision must move into the endpoint stack.

### 3.2 File-delta and low-bandwidth file systems

rsync and the Low-Bandwidth File System (LBFS) show the value of transferring only the differences between file objects that both sides can name and compare [5]. They are the right tool when the workload is genuinely a file tree and both endpoints can run a delta protocol. As our negative results in Section 7.1 confirm, ordinary source-tree updates are better served by rsync than by stream-level reference substitution. ReduLink is not aimed at this case; it is aimed at object or stream delivery where a file-tree delta negotiation is not the serving abstraction.

### 3.3 HTTP delta and shared-dictionary transport

The web has its own lineage of receiver-side reuse over an encrypted channel, and it is the closest existing work. RFC 3229 defines delta encoding in HTTP, sending only the difference of a response relative to a prior instance the client holds, typically using the VCDIFF format [19,20]. Shared Dictionary Compression over HTTP (SDCH) generalized this to a shared dictionary referenced across responses [21], and the recent Compression Dictionary Transport standard (RFC 9842) lets a server designate a previously fetched response as an external Brotli or Zstandard dictionary for future responses [22]. Because these mechanisms operate on HTTP responses, they already run over HTTP/3 on QUIC: receiver-held bytes are reused as a dictionary across an encrypted transport, which is precisely ReduLink's setting.

ReduLink differs in three ways that matter for the deployments it targets. First, it authenticates each reference individually and binds it to epoch, scope, stream, reconstructed offset, length, and a nonce, so a reference cannot be replayed across context or silently expanded; compression-dictionary schemes treat the dictionary as a codec input and rely on the surrounding TLS channel for integrity. Second, ReduLink defines an explicit fail-closed semantic-repair path (MISS then FULL) rather than degrading to a codec fallback. Third, it is a general stream-representation abstraction rather than a codec scoped to HTTP responses with URL-pattern dictionary matching. We therefore treat compression dictionary transport as the strongest deployed point of comparison and position ReduLink as the authenticated, fail-closed, transport-general variant of the same receiver-side reuse idea; Section 12 and Table 25 consolidate this comparison. The distinction is not merely cosmetic: consider a shared or partially trusted origin dictionary, where a CDN edge or several tenants serve from one warm dictionary. An unauthenticated codec reference cannot distinguish a legitimate cross-version reference from a forged or cross-context one, whereas ReduLink's per-reference binding to epoch, scope, stream, offset, and nonce makes such a reference fail closed. That shared-dictionary integrity case, not raw compression ratio, is the gap ReduLink targets.

### 3.4 Content-defined chunking and deduplication privacy

Content-defined chunking determines where chunk boundaries fall and strongly affects deduplication ratio and throughput; recent surveys and fast designs analyze the trade-offs in depth [11,12]. Chunking and deduplication also create content-existence side channels: deduplication can leak whether a receiver already holds particular content, and chunk-boundary parameters can themselves be attacked [13,16,17]. These results directly motivate ReduLink's privacy scoping (Section 4) and its exclusion of global cross-user dictionaries. Modern Ethernet line rates continue to increase [14], which is why the paper avoids physical-rate claims and reports only representation-layer savings.

### 3.5 Deployment model

The deployment model has three roles: a sender with access to the new byte stream, a receiver with a scoped warm dictionary, and a policy domain that defines whether dictionary state is per connection, per origin, or per tenant. Public Internet mode defaults to per-connection dictionaries. Public artifact mode may use signed or manifest-controlled per-origin dictionaries. Enterprise mode may use per-tenant dictionaries with audit and quota controls. Global private cross-user dictionaries remain out of scope. Table 1 summarizes when ReduLink is and is not the right tool relative to rsync and compression.

**Table 1. ReduLink deployment modes and dictionary scoping.**

| Deployment              | Dictionary scope         | Why ReduLink rather than rsync/compression     |
|-------------------------|--------------------------|------------------------------------------------|
| CDN or registry objects | Per origin or per client | Objects arrive as encrypted streams; file-tree |

|                             |                        |                                                                                                              |
|-----------------------------|------------------------|--------------------------------------------------------------------------------------------------------------|
|                             |                        | delta negotiation may not be the serving abstraction.                                                        |
| Backup/page streams         | Per endpoint or tenant | Page-like chunks recur across snapshots and can be referenced while preserving authenticated reconstruction. |
| Enterprise VPN/admin domain | Per tenant             | Policy can allow shared warm state while keeping middleboxes away from plaintext.                            |
| Source-tree sync            | File tree              | Usually better served by rsync; ReduLink is not targeted here.                                               |

## 4. Protocol and Security Model

ReduLink is a representation layer for cooperating, mutually authenticated endpoints. QUIC packetization, packet numbers, ACK processing, TLS, and congestion control remain QUIC responsibilities; ReduLink defines only how application bytes are represented as authenticated FULL, REF, MISS, and DICT\_ACK records and how the receiver validates them. The sender chunks outgoing bytes and emits FULL frames for new chunks or REF frames for chunks already believed present at the receiver. A REF is deliverable only if dictionary membership, epoch, scope, stream id, reconstructed offset, length, nonce, chunk identity, and authentication tag all validate; otherwise the receiver fails closed and requests semantic FULL repair. The full sender/receiver state machines, dictionary-admission rules, and negotiation parameters are specified in the artifact's protocol appendix.

### ReduLink: authenticated reference substitution over an encrypted QUIC stream




Figure 1. ReduLink protocol and deployment architecture. The sender chunks outgoing bytes and emits FULL or authenticated REF records carried as compact binary application data inside an encrypted QUIC stream; the receiver validates each reference against dictionary membership and bound context before byte-exact reconstruction, failing closed and requesting semantic FULL repair on any mismatch.

### 4.1 Frame types and wire model

ReduLink uses four logical frame types. A FULL frame carries chunk bytes and admits them to the receiver dictionary after validation. A REF frame carries an authenticated reference to a chunk already present at the receiver. A MISS frame is a receiver request for FULL fallback when a reference cannot be resolved. A DICT\_ACK frame optionally acknowledges that a receiver admitted a chunk, letting a conservative sender reference only acknowledged chunks. Critically, the `stream_offset` field always denotes the offset in the original reconstructed application byte stream, never an offset in encoded frame bytes. Table 2 lists the bound fields. The offline model charges 24 bytes of framing per FULL and 32 bytes per REF. These modeled overheads are optimistic, not conservative: the artifact's own compact binary wire format costs a measured 108 bytes of metadata per frame (4-byte length prefix, 74-byte fixed header,

scope string, cid and tag) - a 32-byte REF budget cannot even hold the 16-byte chunk identifier plus 16-byte tag that Section 4.5 requires. Section 7.6 re-prices every headline result at the measured cost; a production varint encoding would land between the two.

**Table 2. ReduLink frame types and their authenticated fields.**

| Frame    | Purpose                                   | Bound / carried fields                                                                |
|----------|-------------------------------------------|---------------------------------------------------------------------------------------|
| FULL     | Carry and admit new chunk bytes           | epoch, stream id, stream offset, chunk length, chunk id, payload, auth tag            |
| REF      | Reference a chunk already at the receiver | epoch, stream id, stream offset, original length, chunk id, reference nonce, auth tag |
| MISS     | Request FULL fallback (fail-closed)       | epoch, stream id, stream offset, chunk id                                             |
| DICT_ACK | Acknowledge dictionary admission          | epoch, chunk id, dictionary generation                                                |

## 4.2 Validation and fail-closed reconstruction

The receiver applies a fixed validation order before any byte is delivered, summarized as follows:

- Reject frames whose epoch does not match the active epoch.
- For FULL, validate the authentication tag and chunk identifier, admit the payload to the dictionary, and deliver the bytes at the intended stream offset.
- For REF, validate the authentication tag, expansion bound, stream offset, original length, nonce, and dictionary presence; deliver reconstructed bytes only after all checks succeed.
- Emit MISS when the referenced chunk is unavailable or policy refuses the reference; a REF that cannot become valid in the current epoch triggers an immediate MISS.
- Apply stream ordering and flow control to reconstructed bytes, while applying congestion accounting to transmitted wire bytes.

A reference miss is treated as a synchronization failure, not a corruption event: the receiver delivers no reconstructed bytes, emits MISS, and the sender repairs with a FULL for the same stream id, offset, length, and chunk id. A REF whose original\_length disagrees with the stored chunk length, or that would expand beyond the negotiated bound, is a hard validation failure. Section 11 reports prototypes that exercise exactly these miss, tamper, and replay paths.

## 4.3 Dictionaries, epochs, and key schedule

Dictionaries are scoped per connection by default and per origin or per tenant only under explicit policy; global cross-user dictionaries are out of scope. Implementations use bounded LRU-style eviction, per-origin or per-tenant quotas, and short-lived epochs, and a receiver may refuse admission under memory pressure or policy. In the artifact, chunk identifiers are keyed MACs with domain separation over the active epoch, dictionary scope, and a hash of the chunk bytes (Section 4.5); the protocol specification additionally allows binding stream context and a canonical length encoding where cross-stream reuse is not wanted. In either profile the same bytes under a different epoch or scope yield a different identifier, and length and stream context are always bound in the frame tag even when they are not part of the identifier. Keying differs by artifact path, and we state this precisely: the native aioquic experiment derives its authentication key through the HKDF exporter-style schedule (keyed by an experiment master secret plus ALPN, epoch, scope, connection, and stream context; tests confirm the derived secret changes when any input changes), whereas the standalone secure model and the authenticated-UDP prototype default to a fixed test secret, so their context binding lives in the MAC message, not the key. A production profile keys the schedule from QUIC TLS exporter bytes on every path, which also gives cross-connection freshness: a frame recorded on one connection fails authentication on any other because the key differs.

## 4.4 Threat model

Table 3 maps each security property to the mechanism it requires and to the current artifact status, and Table 4 names the privacy modes and their leakage risks. The artifact implements HMAC binding, chunk-id validation, nonce rejection, length and expansion checks, and fail-closed repair; it does not implement production exporter-derived keys, custom extension-frame parsing, production replay windows, or cross-tenant isolation enforcement, which are stated as future work. Like other deduplication systems, ReduLink can create a content-existence oracle when an adversary can choose payloads or references and observe transfer size, timing, or MISS behavior; per-connection dictionaries

remove the cross-user oracle in public mode, and shared dictionaries are permitted only for public artifacts or explicit trust domains. The security evidence in this paper is artifact-level functional testing plus the reduction-style analysis of Section 4.5; no side-channel mitigations (padding, constant-shape errors, timing equalization) are implemented in the artifact. Dictionary eviction is itself an oracle in shared-dictionary modes: an adversary who can insert FULL traffic can cascade-evict a victim's warm entries (Section 8.2 demonstrates the mechanics under the LRU budget) and then probe REF-versus-MISS behavior to infer content existence, which is why per-connection scoping remains the default and shared dictionaries require explicit policy.

**Table 3. Security properties, required mechanisms, and artifact status.**

| Property          | Required mechanism                                                     | Artifact status                                                          |
|-------------------|------------------------------------------------------------------------|--------------------------------------------------------------------------|
| Integrity         | FULL/REF authentication, chunk-id, offset and length binding           | Modeled by reconstruction and mismatch tests; production crypto pending. |
| Context binding   | Exporter-derived keys, epoch/stream/offset/scope, nonce, replay window | HMAC binding and nonce rejection implemented; exporter keys pending.     |
| Dictionary safety | Authenticated FULL, manifest commitment, admission/eviction policy     | FULL chunk-id checks modeled; manifest policy not implemented.           |
| Expansion bound   | Per-frame, per-stream, per-epoch reconstructed-byte caps               | Length and accounting checks modeled; full flow control pending.         |
| Privacy scope     | Per-connection default, no global cross-user dictionary                | Policy specified; cross-tenant enforcement not implemented.              |

**Table 4. Privacy modes and leakage risks.**

| Mode               | Dictionary scope              | Leakage risk                            | Default           |
|--------------------|-------------------------------|-----------------------------------------|-------------------|
| Public Internet    | Per connection only           | Same-connection access-pattern leakage  | Yes               |
| Public artifacts   | Per-origin signed manifest    | Artifact-version / popularity inference | Optional          |
| Enterprise VPN     | Tenant / admin domain         | Intra-tenant content-existence leakage  | Optional          |
| CDN/update channel | Origin-scoped public versions | Version-possession inference            | Optional          |
| Global cross-user  | Any user                      | Private content-possession leakage      | No (out of scope) |

## 4.5 Formal security model and analysis

We analyze ReduLink's reference authentication against a network adversary in the cooperating-endpoint setting. The endpoints are mutually authenticated by QUIC/TLS and share a per-connection ReduLink key  $K$  derived from the QUIC TLS exporter via HKDF [24] with domain separation over ALPN, epoch, scope, and connection and stream context; the adversary does not know  $K$ . Following the Dolev-Yao model, the adversary observes the encrypted path and may drop, reorder, inject, duplicate, and replay ReduLink records, and in shared-dictionary modes may submit chosen FULL content into a shared dictionary. The adversary's goal is to make a receiver deliver application bytes the legitimate sender did not authenticate for the current context, or to violate the replay or expansion bounds.

We require four properties. Reference unforgeability: no efficient adversary can produce a FULL or REF frame that the receiver accepts and that delivers bytes the sender did not authenticate for the bound context, except with negligible probability. Context binding and replay resistance: a frame authenticated for a given epoch, scope, stream, and offset is rejected in any other context, and a nonce already seen within the epoch window is rejected. Expansion bound: an accepted REF delivers exactly the authenticated chunk length, so a small reference cannot reconstruct unbounded output. Dictionary safety: only authenticated FULL chunks, or signed-manifest chunks, are admitted to the dictionary. The theorem below covers the first two properties, with the replay and expansion bounds following from the same verification checks; manifest-mode dictionary safety is a specification requirement that is not analyzed here because the manifest mechanism is unimplemented (Table 3).

Concretely (Section 4.1), each chunk identifier is a truncated keyed MAC,  $\text{cid} = \text{HMAC}_K(\text{"cid"}, \text{epoch}, \text{scope}, \text{SHA-256}(\text{chunk}))$ , and each frame carries  $\text{tag} = \text{HMAC}_K(\text{kind}, \text{epoch}, \text{scope}, \text{stream}, \text{offset}, \text{length}, \text{nonce}, \text{cid}, \text{SHA-256}(\text{payload}))$  [23]. Verification is authentication-first: the receiver recomputes the tag over the frame's own fields

and rejects on mismatch with a single generic error, so an attacker without the key cannot distinguish which field was tampered with; only authentically-tagged frames proceed to the context checks (epoch, scope, stream, offset) and then to replay rejection against a bounded, reorder-tolerant nonce window. In the artifact the window holds the most recent nonces and treats anything at or below its floor as replayed, bounding receiver memory for long-lived sessions.

Theorem (informal). If HMAC-SHA-256 is existentially unforgeable under chosen-message attack and SHA-256 is collision-resistant, then ReduLink reference substitution satisfies reference unforgeability and context binding in the adversary model above. Proof sketch. Suppose an efficient adversary makes a receiver accept a frame delivering unauthorized bytes for some context  $c$ . Acceptance requires a valid tag over  $c$ ; because the context fields are inside the MAC input and  $K$  is secret, a frame accepted for a context the sender never authenticated yields a valid MAC on a fresh message, contradicting existential unforgeability. The remaining case is a REF that verifies but resolves to bytes other than those referenced, which requires two distinct chunks sharing a  $cid$ , i.e. a collision either in SHA-256 or in the 128-bit truncated identifier output; the former contradicts collision resistance and the latter is birthday-bounded as quantified below. Replay across context fails because the context is bound in the tag and checked at verification; replay within a session fails because the nonce window rejects both repeated nonces and nonces below its floor. One caveat is inherited from the artifact keying: nonces restart at one for each encoding, so replaying an entire recorded session to a fresh receiver holding the same fixed test key would verify - cross-session freshness comes from per-connection derived keys (Section 4.3), which the aioquic path implements and the standalone model does not. A REF whose length disagrees with the stored chunk is additionally caught by the explicit length binding even when identifiers collide. The expansion bound holds because the receiver delivers exactly length bytes from the identified dictionary entry.

The MAC input is a canonical JSON encoding of the labeled fields (sorted keys, fixed separators); an interoperating implementation must reproduce this byte-exact canonicalization or define its own. Two honest caveats apply. First, the artifact truncates both the identifier and the tag to 128 bits, so identifier-collision resistance is birthday-bounded at about  $2^{64}$  and a production profile should size the identifier to the dictionary scale: at  $2^{30}$  dictionary entries the accidental-collision probability with 128-bit identifiers is about  $2^{-69}$ , and even at  $2^{40}$  entries about  $2^{-49}$ , so 128 bits suffices for realistic dictionaries, while deployments aggregating very large multi-tenant dictionaries across many epochs may prefer 192- or 256-bit identifiers for proof margin. The 128-bit tag gives adequate forgery resistance since the adversary lacks the key. Second, the artifact keys HMAC from an explicit test secret rather than the live QUIC exporter, so the analysis describes the production keying profile; the secure-model and authenticated-UDP tests (Section 11) reject tampered tags and replayed nonces, which is consistent with, but does not constitute, a machine-checked proof.

## 5. Implementation and Metrics

The artifact is implemented in Python for reproducibility and uses aioquic for native QUIC stream experiments [18]. ReduLink messages are compact binary application-stream records inside QUIC streams; this exercises a real QUIC handshake and encryption but does not implement custom QUIC frame types or transport parameters. The JSON encoding is retained only as a readable baseline, and the default path uses the compact binary format, which encodes frame type, stream id, reconstructed offset, length, chunk id, nonce, authentication tag, and payload only where required.

Three byte layers are reported. Input bytes are the reconstructed application bytes. Stream-payload bytes are the ReduLink or raw application bytes written into QUIC streams. UDP-estimated bytes add a local IPv4/UDP estimate from the loopback proxy path. The paper avoids treating stream-payload multipliers as full wire-rate measurements; they isolate the representation layer so that raw QUIC, ReduLink, gzip, zstd, fixed reuse, and rsync-style baselines can be compared consistently. Table 5 records, for each evidence layer, what it supports and what it does not prove.

Two reproducibility notes apply to the absolute numbers. Byte multipliers are deterministic functions of the input bytes and chunking parameters and are therefore machine-independent; they are the paper's results. Wall-clock times and local throughputs (for example the client elapsed times in Section 8 and the MiB/s figures in Section 9) are indicative measurements from a single commodity multi-core Linux host and are hardware-dependent; they should be read as order-of-magnitude component costs, not as portable performance guarantees.

**Table 5. Evidence hierarchy: what each layer supports and does not prove.**

| Evidence          | Supports                                                       | Does not prove                         |
|-------------------|----------------------------------------------------------------|----------------------------------------|
| Model             | FULL/REF accounting and reconstruction                         | Transport behavior                     |
| Secure model      | HMAC binding and replay checks                                 | Live QUIC exporter use                 |
| aiquic stream     | Encrypted QUIC stream mapping                                  | Custom extension frames                |
| UDP/IPv4 estimate | Local datagram accounting                                      | Full packet capture                    |
| Workloads         | Workload sensitivity                                           | Universal acceleration                 |
| Path emulation    | Measured shared-bottleneck completion and queueing (userspace) | Kernel netem or Internet-path fairness |

## 6. Evaluation Methodology

Evaluation uses four classes of evidence. First, deterministic journal fixtures (disk snapshot, OCI layer, package metadata, repository snapshot, structured logs, and a compressed negative control) isolate predicted positive and negative workload shapes. Second, public source-release pairs (Click, Redis, nginx) test ordinary source-tree updates where ReduLink should not be assumed to help. Third, object-aligned public release experiments use the same public bytes but model registry/CDN object delivery rather than tarball synchronization. Fourth, native aioquic stream experiments measure the encrypted stream mapping on positive and negative cases, under loss, at scale, and against competing flows and emulated bottlenecks. The journal fixtures are constructed with a chosen unchanged fraction and therefore illustrate workload shapes rather than estimate real-world gains; the only fully external inputs are the public source releases, which are negative for ReduLink, and their object-aligned re-framing, which is the paper's primary external positive signal. An author-independent-overlap version-pair study of real package upgrades is added in Section 7.5, and Section 7.6 re-prices all headline results at the measured wire framing and adds a dictionary-delta baseline.

For each workload the paper reports ReduLink fixed chunking, ReduLink content-defined chunking where relevant, fixed-block reuse, compression, and rsync where the baseline is structurally applicable. The fixed-block baseline is intentionally strong and simple: if it beats ReduLink, the result is reported rather than hidden, because ReduLink is a security and transport-compatibility layer over reuse, not a superior matching algorithm. All external corpora are hash-pinned, and every figure and table is regenerated from committed result files.

## 7. Workload Results and Baseline Interpretation

Table 6 reports the deterministic warm-state fixtures. On most byte-stable workloads, fixed-block reuse matches or beats ReduLink; the distinct ReduLink contribution is not superior matching but authenticated, scoped, stream-compatible reference substitution with explicit repair and privacy rules. The negative rows (package metadata, logs, and the compressed control) are included precisely because they fail: where repetition is local to one object, compression wins, and where there is no reusable structure, both fixed reuse and ReduLink correctly decline to help. The Unchanged column makes the mechanism explicit: because for aligned whole-chunk reuse the effective multiplier is bounded by roughly one over one minus the unchanged fraction (the fixed-block-reuse column can exceed this bound because its byte-scan alignment also matches moved content), these deterministic fixtures (96 to 100 percent unchanged on the positive rows) are illustrative workload shapes chosen to expose where reference substitution can and cannot help, not estimates of real-world overlap. The repository fixture in particular changes only 210 of 501,760 bytes, so its 24.73x is close to a no-op transfer.

**Table 6. Warm-state workload results (effective stream-payload multiplier), with the constructed unchanged fraction of each fixture.**

| Workload        | Unchanged | ReduLink | Fixed reuse | Best compression | RL - fixed |
|-----------------|-----------|----------|-------------|------------------|------------|
| disk snap       | 97.3%     | 28.49x   | 31.97x      | 1.06x            | -3.48x     |
| oci layer       | 96.5%     | 11.05x   | 10.68x      | 1.50x            | +0.37x     |
| package meta    | 3.5%      | 0.99x    | 9.88x       | 14.31x           | -8.89x     |
| repository snap | 100.0%    | 24.73x   | 24.60x      | 10.37x           | +0.14x     |
| logs            | 3.3%      | 0.99x    | 14.89x      | 18.01x           | -13.89x    |



|                |      |       |       |       |        |
|----------------|------|-------|-------|-------|--------|
| compressed neg | 0.4% | 0.99x | 1.00x | 1.00x | -0.01x |
|----------------|------|-------|-------|-------|--------|

## 7.1 External Public Source Releases

Table 7 reports three independent, hash-pinned source-release pairs. The result is negative for ReduLink and strongly positive for rsync. This is an important and deliberately reported finding: ordinary related source trees are better served by file-oriented delta transfer than by stream-level reference substitution, because rsync can exploit fine-grained intra-file similarity that whole-object referencing cannot.

**Table 7. External public source-release pairs: ReduLink and fixed reuse versus real rsync (negative case for ReduLink).**

| Public pair            | ReduLink | Fixed reuse | rsync total |
|------------------------|----------|-------------|-------------|
| click-8.1.7 to 8.1.8   | 0.99x    | 0.99x       | 6.48x       |
| redis-7.2.4 to 7.2.5   | 1.00x    | 1.00x       | 73.13x      |
| nginx-1.25.3 to 1.25.4 | 1.00x    | 0.99x       | 49.49x      |

## 7.2 Object-Aligned Public Release Transfers

Table 8 and Figure 2 use the same public release tarballs but change the transfer abstraction: files are treated as object records and chunked independently, modeling registry/CDN object delivery rather than source-tree synchronization. Redis and nginx show positive ReduLink results because many public release objects remain byte-identical across versions; Click is weaker because fewer files are unchanged. ReduLink tracks the idealized fixed-object-reuse baseline within a few percent, and the authenticated variant costs a further 3 to 11 percent (11.3 percent on the layer-like case, 6.8 percent on redis). This is a stronger external positive signal than a purely synthetic corpus, though it is still a transfer-model experiment rather than a captured production trace.

**Table 8. Object-aligned public release transfers modeling registry/CDN object delivery.**

| Public release | File stability      | Bytes      | ReduLink | Secure RL | Fixed object reuse | gzip  |
|----------------|---------------------|------------|----------|-----------|--------------------|-------|
| click          | 80/144 unchanged    | 947,826    | 1.94x    | 1.91x     | 1.94x              | 2.86x |
| redis          | 1557/1582 unchanged | 15,467,095 | 9.70x    | 9.04x     | 9.68x              | 4.61x |
| nginx          | 468/503 unchanged   | 7,357,392  | 4.38x    | 4.25x     | 4.38x              | 6.09x |




Figure 2. Object-aligned public release transfers: effective stream-payload multipliers for ReduLink, the authenticated variant, fixed object reuse, and gzip.

## 7.3 Layer-Like External-Positive Case

Table 9 reports a Redis-derived layer-like case built from included public Redis release bytes and aligned changed blocks. Its role is to test the predicted positive case for layer-like objects where stable chunks remain aligned; it is not a production trace. Here ReduLink reaches 21.37x with the authenticated variant at 18.96x, again close to the fixed-reuse reference.

**Table 9. Layer-like external-positive case derived from public Redis release bytes.**

| Case                          | Input bytes | ReduLink | Secure RL | Fixed reuse |
|-------------------------------|-------------|----------|-----------|-------------|
| redis-layered-public-positive | 524,288     | 21.37x   | 18.96x    | 21.33x      |

## 7.4 Block-Size Sensitivity

Table 10 and Figure 3 show that the best block size is workload-dependent: smaller blocks capture more aligned reuse but pay more per-chunk overhead, while larger blocks lose alignment. The artifact therefore treats 4 KiB as a reproducible default rather than a universal optimum, and a production deployment would tune the chunk size per workload class. The content-defined chunker consistently underperforms fixed chunking on these aligned fixtures (at 4 KiB: 10.71x versus 28.49x on the disk snapshot, 4.60x versus 11.05x on the OCI layer), because the fixtures change bytes in place without insertions; CDC's insertion robustness is not exercised here, and its Python implementation is also the slowest component in Table 17.

**Table 10. Block-size sensitivity of the effective multiplier (fixed chunking).**

| Workload            | 1 KiB  | 2 KiB  | 4 KiB  | 8 KiB  | 16 KiB |
|---------------------|--------|--------|--------|--------|--------|
| disk snapshot       | 17.13x | 23.33x | 28.49x | 17.08x | 8.99x  |
| oci layer           | 12.11x | 11.86x | 11.05x | 5.91x  | 3.60x  |
| repository snapshot | 20.21x | 24.99x | 24.73x | 14.45x | 7.55x  |




Figure 3. Block-size sensitivity of the effective multiplier for fixed chunking across three warm-update fixtures.

## 7.5 Package version-pair study (PyPI)

The fixtures of Section 7 and the object-aligned experiment use bytes the author selected; to obtain a positive case from fully independent data, Table 11 reports a version-pair study over real released packages. For six widely used Python packages we downloaded the wheel distributions of two consecutive released versions directly from PyPI,

recording the SHA-256 of every wheel for reproducibility, treated each file inside a wheel as an object a client already holds after installing the older version, and measured object reuse with the same accounting as Section 7.2; reconstruction was additionally confirmed byte-exact by a real model round trip. The result is honestly mixed and matches the claim boundary: patch releases with high file stability (rich, 76 of 82 files unchanged; click) reach about 11x to 12x, whereas feature or minor-version upgrades (jinja2, urllib3, packaging) change enough files that gzip wins and ReduLink falls to roughly 1.1x to 1.9x. This is the paper's strongest external positive evidence because the byte overlap is whatever the real releases happened to share; the package list and version pairs were still author-selected, so this is a version-pair study rather than a captured client trace, and no claim about upgrade-traffic frequency or population weighting is made. Two modeling caveats apply. PyPI today serves each release as a single compressed wheel archive, so this experiment models a file-granular registry object channel, the deployment class of Section 7.2, rather than pip's current wire format; realizing these gains in practice would require registry-side per-object serving. For the same reason the gzip baseline is computed over the uncompressed object stream, whereas production wheels are already deflate-compressed in transit, so the comparison is internally consistent across methods but should not be mapped directly onto today's PyPI traffic.

**Table 11. PyPI version-pair study: real package upgrades (object-aligned, hash-pinned).**

| Package   | Versions       | Files unch. | ReduLink | Secure RL | Fixed reuse | gzip  | Recon. |
|-----------|----------------|-------------|----------|-----------|-------------|-------|--------|
| rich      | 13.7.0->13.7.1 | 76/82       | 12.41x   | 11.40x    | 12.39x      | 4.43x | OK     |
| jinja2    | 3.1.3->3.1.4   | 7/30        | 1.10x    | 1.09x     | 1.09x       | 4.03x | OK     |
| click     | 8.1.6->8.1.7   | 13/21       | 11.17x   | 10.40x    | 11.14x      | 3.98x | OK     |
| werkzeug  | 3.0.1->3.0.2   | 48/69       | 3.06x    | 2.99x     | 3.05x       | 3.97x | OK     |
| urllib3   | 2.1.0->2.2.0   | 19/39       | 1.87x    | 1.84x     | 1.87x       | 3.81x | OK     |
| packaging | 23.2->24.0     | 7/20        | 1.54x    | 1.52x     | 1.53x       | 3.79x | OK     |

## 7.6 Framing sensitivity and a dictionary-delta baseline

Two review-driven checks close out the workload evaluation; Table 12 reports both. First, framing repricing: the offline model's 24/32-byte overheads understate the artifact's own compact binary wire format, which costs a measured 108 bytes of metadata per frame. Re-pricing every FULL and REF at the measured cost lowers the headline multipliers by 5 to 28 percent - redis falls from 9.70x to 7.90x, the layer-like case from 21.37x to 15.28x - without changing any qualitative conclusion; both figures are reported so readers can price a production variant encoding anywhere between them. Second, the strongest byte baseline we know of: zstd with the prior object stream as a raw-content dictionary (zstd -3 --patch-from [26]), the deployed-style analog of Compression Dictionary Transport. Its result is unambiguous and we state it plainly: whenever the receiver retains the exact prior byte stream and a codec delta is acceptable, dictionary delta dominates ReduLink on bytes by one to three orders of magnitude (65x to 1,785x on these corpora). ReduLink's justification is therefore never byte-optimality against a whole-stream codec delta; it is per-reference authentication, object-granular incremental delivery against a chunk dictionary rather than a contiguous retained stream, fail-closed repair, and stream compatibility. Deployments whose receivers can hold the exact prior stream and tolerate codec-level trust should prefer dictionary delta; Section 12 revisits this boundary.

**Table 12. Framing repricing (measured 108-byte wire framing vs 24/32-byte model) and the zstd dictionary-delta baseline on the same object streams.**

| Case                          | Input bytes | RL (model) | RL (wire-priced) | zstd --patch-from | gzip  |
|-------------------------------|-------------|------------|------------------|-------------------|-------|
| click-8.1.7-to-8.1.8          | 947,826     | 1.94x      | 1.84x            | 65.25x            | 2.86x |
| redis-7.2.4-to-7.2.5          | 15,467,095  | 9.70x      | 7.90x            | 1785.43x          | 4.61x |
| nginx-1.25.3-to-1.25.4        | 7,357,392   | 4.38x      | 4.00x            | 951.72x           | 6.09x |
| redis-layered-public-positive | 524,288     | 21.37x     | 15.28x           | 1659.14x          | 4.24x |
| pypi rich-13.7.0-to-13.7.1    | 947,654     | 12.41x     | 9.72x            | 603.15x           | 4.43x |
| pypi jinja2-3.1.3-to-3.1.4    | 491,214     | 1.10x      | 1.07x            | 156.44x           | 4.03x |
| pypi click-8.1.6-to-8.1.7     | 353,767     | 11.17x     | 9.08x            | 639.80x           | 3.98x |
| pypi werkzeug-3.0.1-to-3.0.2  | 762,091     | 3.06x      | 2.86x            | 387.09x           | 3.97x |
| pypi urllib3-2.1.0-to-2.2.0   | 391,321     | 1.87x      | 1.79x            | 29.49x            | 3.81x |
| pypi packaging-23.2-to-24.0   | 174,172     | 1.54x      | 1.48x            | 82.47x            | 3.79x |

## 8. Native QUIC Experiments

### 8.1 Stream mapping under loss

Table 13 compares raw and ReduLink stream mapping at zero and periodic datagram loss, and Table 14 reports positive and negative workload cases. ReduLink reconstructs byte-exactly in every case, including the lossy proxy path, and the compressed-negative control correctly yields no gain. Figure 4 plots the workload cases. Because ReduLink writes fewer stream-payload bytes, its UDP-estimated multiplier stays above one on positive workloads even after the IPv4/UDP per-datagram overhead is added. One apparent inconsistency deserves explicit reconciliation: the Redis-layered corpus reaches 21.37x in the offline model (Table 9) but 3.83x over the QUIC mapping (Table 13). The gap decomposes into three deliberate differences: the QUIC experiment chunks at 1 KiB rather than 4 KiB (four times as many frames), each frame pays the measured 108-byte wire framing rather than the 32-byte model charge (Section 7.6), and the server dictionary is deliberately punctured (every seventh chunk withheld), forcing 72 semantic misses whose FULL repairs carry payload. No offline multiplier in Tables 6 to 12 has been achieved over the wire; the QUIC tables are the end-to-end evidence, and the offline tables isolate the representation layer.

**Table 13. Native aioquic stream mapping: raw versus ReduLink under zero and periodic datagram loss.**

| Method   | Loss every | Stream x | UDP-est x | Recon. |
|----------|------------|----------|-----------|--------|
| Raw      | 0          | 1.00x    | 0.90x     | OK     |
| ReduLink | 0          | 3.19x    | 2.60x     | OK     |
| Raw      | 9          | 1.00x    | 0.78x     | OK     |
| ReduLink | 9          | 3.19x    | 2.38x     | OK     |

**Table 14. Native aioquic stream mapping across positive and negative workload cases.**

| Case                | Input   | Stream x | UDP-est x | Misses | Recon. |
|---------------------|---------|----------|-----------|--------|--------|
| demo positive       | 98,304  | 3.19x    | 2.57x     | 13     | OK     |
| compressed negative | 262,242 | 0.91x    | 0.84x     | 0      | OK     |
| Redis layered       | 524,288 | 3.83x    | 3.46x     | 72     | OK     |




Figure 4. Native aioquic stream mapping on positive and negative workload cases.

## 8.2 Repeated trials and scaling

Table 15 reports three repeated trials of the native demo workload. The multiplier is near-deterministic for fixed bytes (encoded bytes vary by at most one byte across rounds from control-message framing), while elapsed times vary modestly across runs; reporting both avoids cherry-picking a single favorable QUIC run; with only three trials these are run-to-run stability indicators rather than a statistically powered variance estimate. Table 16 then scales the payload from 96 to 16,384 blocks (96 KiB to 16 MiB): the stream multiplier rises from 3.19x to 3.94x as more reusable chunks accumulate, and reconstruction remains byte-exact at every scale. The 16 MiB rows expose an operating constraint that smaller runs cannot show: with the artifact's default dictionary budget of 8,192 chunks, a 16,384-chunk warm set overflows the LRU, sequential FULL admissions cascade-evict the remaining warm entries, and ReduLink degrades to a FULL-only stream (0.92x, pure framing overhead) while still reconstructing byte-exactly, i.e. it fails safe rather than wrong. Raising the budget to 24,576 chunks restores 3.96x and, in this single-run comparison, roughly halves completion time (592 versus 1,305 ms). Dictionary sizing relative to the warm working set is therefore a hard provisioning requirement, not a tuning nicety.

**Table 15. Repeated native QUIC trials (n = 3): run-to-run stability of multipliers and elapsed time.**

| Metric                      | n  | Mean       | Std. dev. | Min       | Max        |
|-----------------------------|----|------------|-----------|-----------|------------|
| raw client ms               | 20 | 101.776500 | 2.542158  | 98.555000 | 107.705000 |
| redulink stream multiplier  | 20 | 3.190036   | 0.000023  | 3.190031  | 3.190135   |
| redulink udp est multiplier | 20 | 2.605427   | 0.003688  | 2.596856  | 2.608571   |
| redulink client ms          | 20 | 99.743800  | 0.972249  | 98.298000 | 101.618000 |

**Table 16. Payload scaling of the native aioquic stream mapping, including the dictionary-budget overflow and recovery at 16 MiB.**

| Blocks | Input bytes | Dict budget | Stream x | Misses | Client ms | Recon. |
|--------|-------------|-------------|----------|--------|-----------|--------|
| 96     | 98,304      | 8,192       | 3.19x    | 13     | 98.4      | OK     |
| 512    | 524,288     | 8,192       | 3.77x    | 73     | 123.7     | OK     |
| 1024   | 1,048,576   | 8,192       | 3.87x    | 146    | 145.8     | OK     |
| 4096   | 4,194,304   | 8,192       | 3.94x    | 585    | 229.1     | OK     |
| 16384  | 16,777,216  | 8,192       | 0.92x    | 0      | 1305.3    | OK     |
| 16384  | 16,777,216  | 24,576      | 3.96x    | 2340   | 592.0     | OK     |

## 9. Component Costs

Table 17 reports component-level throughput on the reference host. Fixed chunking, HMAC validation, and compact binary encoding are fast; the Python content-defined chunker is slow and is explicitly not production line-rate evidence. A production implementation would require optimized native chunking and tighter integration with the transport stack. As noted in Section 5, these throughputs are hardware-dependent component costs, not portable guarantees.

**Table 17. Component-level cost measurements (local Python artifact timing; not production line-rate).**

| Component             | Input bytes | MiB/s local | Roundtrip      |
|-----------------------|-------------|-------------|----------------|
| fixed chunking        | 2,097,152   | 6237.0      | not applicable |
| CDC chunking          | 2,097,152   | 7.5         | not applicable |
| model fixed roundtrip | 1,573,964   | 680.0       | OK             |
| model CDC roundtrip   | 1,573,964   | 3.4         | OK             |
| HMAC roundtrip        | 1,573,964   | 192.2       | OK             |
| HMAC decode           | 1,573,964   | 502.9       | not applicable |
| wire encode           | 1,573,964   | 4696.9      | not applicable |
| wire decode           | 1,573,964   | 2451.2      | not applicable |

## 10. Fairness and Accounting

The central fairness rule is that ReduLink must not receive congestion credit for reconstructed bytes: congestion control and bottleneck service account encoded wire bytes only. Table 18 reports the wire-byte accounting experiment, in which ReduLink's encoded wire share is 0.083 against the raw stream's 0.917 over 16 rounds, while still reconstructing the same application bytes; the effective application multiplier on that 48 KiB wire-accounting fixture is 10.72x. This confirms that the savings appear as fewer encoded bytes, which is the only place a fair transport should credit them.

**Table 18. Wire-byte fairness accounting: ReduLink is charged on encoded bytes, not reconstructed bytes.**

| Metric                           | Value                                                      |
|----------------------------------|------------------------------------------------------------|
| Fairness rule                    | Congestion and bottleneck service count encoded wire bytes |
| Raw encoded wire share           | 0.917                                                      |
| ReduLink encoded wire share      | 0.083                                                      |
| ReduLink uses less wire than raw | yes                                                        |
| Effective application multiplier | 10.72x                                                     |
| Rounds                           | 16                                                         |

Table 19 reports a measured concurrent competing-flow experiment in which a raw QUIC flow and a ReduLink flow (both using aioquic's default NewReno congestion controller, so fairness is between identical stacks) run simultaneously over a shared localhost schedule with periodic datagram loss, repeated over 20 rounds. Across the measured rounds the reconstructed-rate Jain index [25] is 0.975 and the encoded-rate index is 0.800 (for two concurrent flows Jain's index ranges from 0.5 to 1.0), with every flow reconstructing correctly; the lower encoded-rate index is expected, reflecting that ReduLink transmits far fewer encoded bytes than the raw flow rather than transport-level unfairness. This is a real loopback measurement without rate shaping; Table 20 adds measured rate and delay emulation, Table 21 varies the semantic miss rate on one constrained point, and Table 22 reports bootstrap confidence intervals for the repeated QUIC measurements. Table 23 is an analytic illustration rather than a measurement: it applies a closed-form fluid model - fair-share serialization of the measured encoded bytes plus one RTT of completion latency - to the measured encoded byte counts. Because ReduLink encodes fewer bytes, the model yields a lower completion time and therefore a higher implied application rate at the same encoded fair share; for example at 25 Mbps and 10 ms the model gives 29.7 ms versus 51.3 ms for raw. This advantage is arithmetic, a direct consequence of sending fewer encoded bytes; the measured emulation of Table 20 shows why it is optimistic, since the fluid model ignores queueing delay and ReduLink's application-level repair round trips, which in measurement reverse the completion-time ordering at 5 Mbps. We retain the model only as a best-case bound on the byte-volume effect.

**Table 19. Measured concurrent competing-flow experiment (Jain fairness index, 20 rounds).**

| Metric                        | Value                                                                  |
|-------------------------------|------------------------------------------------------------------------|
| Scenario                      | 20 concurrent rounds, unshaped localhost (no rate limit), loss every 9 |
| Encoded-rate Jain index       | 0.800                                                                  |
| Reconstructed-rate Jain index | 0.975                                                                  |
| All flows reconstructed       | yes                                                                    |
| Scope                         | localhost measured experiment; not a kernel-netem path-emulation study |

Table 20 adds rate and delay to the same harness as a measured emulation: both flows run concurrently through a full-duplex shared bottleneck (one token bucket per direction plus one-way propagation delay), sweeping 5 and 20 Mbps against 20 and 80 ms RTT with 20 rounds per point. This corrects a methodological error we reported in an earlier version. Our first shaper used a single half-duplex bucket, which forced ReduLink's reverse MISS/repair traffic to serialize behind the raw flow's forward bulk data and produced a spurious always-slower finding. With a correct full-duplex link and a larger 20-round sample, byte-stable demo content is near parity at the most constrained point, faster at two grid points, and near parity again once the link is less congested and RTT-bound. We report this openly because it is the point: emulation topology, not just byte counts, determines completion time.

The completion-time benefit is conditional on the miss rate, which Table 20 also shows by pushing real Redis-layered public bytes through the same path instead of the byte-stable demo. That workload punctures the dictionary heavily (72 semantic misses), so ReduLink's payload-bearing repairs and their reverse round trips can erode the byte savings. In the refreshed 20-round grid, Redis-layered ReduLink is slightly faster at 5 Mbps/20 ms and 20 Mbps/80 ms, but slower at 5 Mbps/80 ms and 20 Mbps/20 ms. Reconstruction is byte-exact in every round. The honest synthesis is therefore neither the promotional fewer-bytes-so-faster nor our earlier artifactual always-slower: on a fair full-duplex link, ReduLink's byte savings can shorten completion time when repair traffic and scheduling align, and can be offset by MISS and repair round trips when the warm dictionary is heavily punctured. Fairness itself is unaffected: both flows use identical NewReno stacks, so ReduLink can never gain congestion credit for reconstructed bytes.

**Table 20. Measured full-duplex shared-bottleneck path emulation: mean completion time (ms, +/- s.d.) for competing raw and ReduLink flows on byte-stable demo bytes and real Redis-layered bytes (20 rounds per scenario).**

| Payload | Rate Mbps | RTT ms | Raw ms       | ReduLink ms  | RL/raw | Jain (enc) | Recon. |
|---------|-----------|--------|--------------|--------------|--------|------------|--------|
| demo    | 5.0       | 20.0   | 971.2±242.4  | 949.1±73.2   | 0.98x  | 0.781      | OK     |
| demo    | 5.0       | 80.0   | 1213.8±53.1  | 1062.4±28.6  | 0.88x  | 0.817      | OK     |
| demo    | 20.0      | 20.0   | 367.4±22.0   | 327.1±21.5   | 0.89x  | 0.813      | OK     |
| demo    | 20.0      | 80.0   | 778.3±10.8   | 793.3±9.6    | 1.02x  | 0.781      | OK     |
| redis   | 5.0       | 20.0   | 2096.2±66.4  | 1986.5±102.4 | 0.95x  | 0.756      | OK     |
| redis   | 5.0       | 80.0   | 3232.1±708.3 | 4086.0±694.8 | 1.26x  | 0.694      | OK     |
| redis   | 20.0      | 20.0   | 1095.6±68.2  | 1224.8±60.8  | 1.12x  | 0.721      | OK     |
| redis   | 20.0      | 80.0   | 1438.3±18.9  | 1176.5±33.4  | 0.82x  | 0.790      | OK     |

Table 21 turns the qualitative miss-rate claim into a sensitivity check. It uses the same native aioquic stream mapping and full-duplex userspace shaper at the constrained 5 Mbps, 20 ms point, but varies receiver dictionary thinning on the byte-stable demo payload. As the semantic miss fraction rises from about 1 percent to about 49 percent, ReduLink's stream multiplier falls from 5.75x to 1.48x and the completion-time advantage narrows from roughly 0.64-0.74x to 0.85x of the raw flow. This does not replace a broader rate/RTT grid, and the first row has high raw-flow timing dispersion, but it directly supports the mechanism-level explanation: repair payloads and reverse MISS/FULL round trips progressively consume the byte-saving advantage while preserving byte-exact reconstruction.

**Table 21. Native QUIC miss-rate sensitivity at 5 Mbps and 20 ms RTT on byte-stable demo bytes (three rounds per point).**

| Dictionary thinning | Miss fraction | Raw ms      | ReduLink ms | RL/raw | Stream x | Recon. |
|---------------------|---------------|-------------|-------------|--------|----------|--------|
| 0                   | 1.1%          | 925.8±289.9 | 664.5±76.3  | 0.74x  | 5.75x    | OK     |
| 16                  | 6.7%          | 1284.5±3.3  | 825.5±5.7   | 0.64x  | 4.31x    | OK     |
| 8                   | 13.3%         | 1319.5±1.8  | 968.2±12.5  | 0.73x  | 3.31x    | OK     |
| 4                   | 26.7%         | 1662.4±11.7 | 1322.3±1.3  | 0.80x  | 2.27x    | OK     |
| 2                   | 48.9%         | 2018.9±11.1 | 1722.8±12.5 | 0.85x  | 1.48x    | OK     |

Table 22 adds uncertainty estimates to the repeated QUIC measurements rather than relying only on point estimates. The intervals are deterministic percentile bootstrap confidence intervals over paired raw/ReduLink rows where available. The 20-round refresh strengthens the byte-accounting interpretation and also narrows one earlier overstatement: at the constrained demo point, ReduLink's encoded-byte ratio is stable near 0.31, but the completion-time interval crosses parity, so the fair claim is byte reduction with roughly neutral completion time on that point. The Redis constrained point shows both a stable encoded-byte reduction and a modest mean completion-time advantage in this run, while other Redis grid points still show that semantic MISS/FULL repair round trips can dominate. macOS pf/dnctl and Linux tc/netem kernel-path harnesses were added to the artifact and their dry-run rules are tested, but no live kernel-path aioquic sweep is claimed here: on this development Mac, raw aioquic handshakes intermittently timed out under loopback dummynet, and Linux tc/netem requires a privileged Linux host.

**Table 22. Statistical uncertainty for repeated native QUIC measurements (deterministic bootstrap 95% confidence intervals).**

| Result                                  | n  | Mean  | 95% CI         | Source                   |
|-----------------------------------------|----|-------|----------------|--------------------------|
| Demo constrained completion ratio       | 20 | 1.012 | [0.948, 1.070] | quic_emulated_path       |
| Demo encoded-byte ratio                 | 20 | 0.313 | [0.313, 0.313] | quic_emulated_path       |
| Redis constrained completion ratio      | 20 | 0.948 | [0.932, 0.963] | quic_emulated_path_redis |
| Redis encoded-byte ratio                | 20 | 0.261 | [0.261, 0.261] | quic_emulated_path_redis |
| Concurrent localhost completion ratio   | 20 | 0.921 | [0.826, 1.031] | quic_competing_flows     |
| Concurrent localhost encoded-byte ratio | 20 | 0.313 | [0.313, 0.313] | quic_competing_flows     |
| Sequential UDP-estimated multiplier     | 20 | 2.605 | [2.604, 2.607] | repeated_quic_trials     |

**Table 23. Analytic fluid-model completion times and implied application rates from measured encoded stream bytes (raw versus ReduLink); computed, not measured.**

| Rate Mbps | RTT ms | Raw compl. ms (model) | RL compl. ms (model) | Raw app rate Mbps (model) | RL app rate Mbps (model) |
|-----------|--------|-----------------------|----------------------|---------------------------|--------------------------|
| 10.0      | 10.0   | 113.3                 | 59.3                 | 6.9                       | 13.3                     |
| 10.0      | 50.0   | 153.3                 | 99.3                 | 5.1                       | 7.9                      |
| 25.0      | 10.0   | 51.3                  | 29.7                 | 15.3                      | 26.5                     |
| 25.0      | 50.0   | 91.3                  | 69.7                 | 8.6                       | 11.3                     |
| 100.0     | 10.0   | 20.3                  | 14.9                 | 38.7                      | 52.7                     |
| 100.0     | 50.0   | 60.3                  | 54.9                 | 13.0                      | 14.3                     |

The package includes reviewer-runnable macOS dummynet and Linux tc/netem harnesses, but does not include a completed kernel-emulator or Mininet congestion-control result. The competing-flow and bottleneck results are therefore deliberately conservative emulations that isolate the encoded-byte accounting principle rather than claim production fairness.

## 11. Repair and Authentication Prototypes

Beyond the in-stream experiments, three prototypes exercise the fail-closed and authentication paths directly. Table 24 summarizes them. The semantic-repair demo forces a 25 percent dictionary-miss fraction and confirms that every miss is repaired with a FULL and that reconstruction is byte-exact. The localhost UDP repair experiment drops every seventh data datagram and uses timeout-based retransmission, again reconstructing fully after 15 semantic repairs and 6 retransmissions. The authenticated UDP experiment adds explicit negative probes: a tampered authentication tag and a replayed nonce are both rejected (1 tamper and 1 replay rejection) before normal authenticated repair traffic is accepted, demonstrating fail-closed behavior under active manipulation.

**Table 24. Repair and authentication prototypes (localhost).**

| Prototype | Input bytes | Misses / repairs | Negative probes rejected | Recon. |
|-----------|-------------|------------------|--------------------------|--------|
|-----------|-------------|------------------|--------------------------|--------|



|                           |        |         |                    |    |
|---------------------------|--------|---------|--------------------|----|
| Semantic repair demo      | 32,546 | 2 / 2   | n/a (25% missing)  | OK |
| UDP repair (drop every 7) | 49,152 | 15 / 15 | 6 retransmits      | OK |
| Authenticated UDP         | 65,536 | 12 / 12 | 1 tamper, 1 replay | OK |

## 12. Why Not Just Compress or Use rsync?

Compression is best when repetition is local to one object; rsync is best when both endpoints share a file tree and can run a delta protocol; HTTP delta or shared-dictionary transport is best when the application is HTTP and a codec-level dictionary suffices; and whole-stream dictionary delta (zstd --patch-from, Section 7.6) is best on raw bytes whenever the receiver retains the exact prior stream - by one to three orders of magnitude on our corpora. The ReduLink use case is different: encrypted endpoint streams where dictionary state is already available at the receiver, the transfer is not naturally a file-tree synchronization session, and each reconstruction step must be individually authenticated and fail closed. The evaluation confirms the boundary. On source-release trees, rsync dominates (Table 7). On package metadata and logs, compression or fixed-block reuse dominates (Table 6). On object-aligned public releases and layer-like byte-stable objects (Tables 8 and 9), ReduLink provides authenticated reference substitution at modest overhead over a simple fixed-reference baseline.

The motivating cases are therefore not arbitrary file synchronization. They are stream-serving cases in which the endpoint already knows that a receiver has dictionary state but must still authenticate every reconstruction step and preserve encrypted transport semantics. A CDN edge, registry, backup client, or enterprise service may choose this representation because it avoids exposing plaintext to a WAN optimizer and avoids changing the application protocol into a file-tree delta session, while keeping a stronger per-reference integrity guarantee than a codec-level shared dictionary provides.

Table 25 consolidates the positioning. ReduLink is the only approach in the comparison that combines operation over an end-to-end encrypted QUIC stream with per-reference authentication and an explicit fail-closed repair path; the other approaches each win in their own regime, which is exactly why ReduLink is presented as a complementary representation layer rather than a replacement.

**Table 25. Positioning of ReduLink relative to related redundancy-suppression approaches.**

| Approach                                  | E2E-encrypted transport      | Per-reference auth       | Fail-closed repair                   | Receiver dictionary         | Best-fit workload                                        |
|-------------------------------------------|------------------------------|--------------------------|--------------------------------------|-----------------------------|----------------------------------------------------------|
| In-network RE [1-4]                       | No (needs plaintext vantage) | No                       | n/a                                  | Network cache               | High-redundancy unencrypted links                        |
| rsync / LBFS [5]                          | Endpoints only               | Transport integrity only | n/a (delta negotiation)              | File/object index           | File-tree synchronization                                |
| HTTP shared dictionary / CDT [19-22]      | Yes (HTTP/3 on QUIC)         | No (TLS channel only)    | No (codec fallback, not fail-closed) | Prior HTTP response         | HTTP responses with prior versions                       |
| Local compression (gzip/zstd)             | Yes                          | No                       | n/a                                  | None (intra-object)         | Locally repetitive single objects                        |
| zstd dictionary delta (--patch-from) [26] | Yes (as payload)             | No (codec trust)         | No                                   | Exact retained prior stream | Full prior stream retained; best byte savings (Sec. 7.6) |
| ReduLink (this work)                      | Yes (QUIC streams)           | Yes (per reference)      | Yes (MISS then FULL)                 | Scoped warm dictionary      | Warm-state object / layer transfers                      |

## 13. Limitations and Future Work

The implementation maps ReduLink into QUIC streams rather than adding custom QUIC extension frames or transport parameters, and it uses an exporter-style key schedule model rather than live QUIC TLS exporter bytes. The public object-sequence and layer-like workloads are derived from real release bytes but are transfer-model experiments rather than captured production traces. The fairness evidence is a measured 20-round competing-flow experiment plus a measured 20-round full-duplex userspace path emulation across a rate/RTT grid on both byte-stable and real repair-heavy payloads (Section 10); userspace shaping is faithful at the studied rates but is not kernel tc/netem or Mininet, and all transport experiments run over localhost. We added macOS pf/dnctl dummynet and Linux tc/netem harnesses and verified their dry-run rules, but we do not report a completed kernel-path aioquic sweep: raw aioquic handshakes timed out intermittently under loopback dummynet on this Mac, and Linux tc/netem must be run on a privileged Linux

host. The path-emulation history is a caution in itself: our first single-bucket shaper produced a completion-time finding that a correct full-duplex model reversed, so the completion-time results should be read as emulation-model-dependent until confirmed on a kernel emulator or real path. The miss-rate sweep is a one-path sensitivity check, not a full transport parameter study. Native-QUIC experiments now reach 16 MiB over a single localhost connection, including a demonstrated dictionary-budget overflow and recovery; multi-connection behavior, larger-than-memory dictionaries, and Internet-path dynamics remain out of scope. The independent positive version-pair study of Section 7.5 captures real package-distribution objects rather than a live production registry pull. The security analysis of Section 4.5 reduces reference unforgeability to standard HMAC and collision-resistance assumptions, but a machine-checked proof and production exporter-derived keying on every path remain future work. Headline byte multipliers are model-framed; the wire-priced values of Section 7.6 are 5 to 28 percent lower, and whole-stream dictionary delta beats both wherever its retention and trust assumptions hold. Absolute timings are hardware-dependent.

Future work should add custom QUIC frame integration and transport-parameter negotiation, live exporter-derived keys, production replay-window and cross-tenant isolation enforcement, a completed privileged-host tc/netem or Mininet fairness study with competing real flows, packet capture, larger and live-captured registry or backup traces, a broader miss-rate grid across bandwidth and RTT, a machine-checked security proof, and a direct head-to-head comparison against Compression Dictionary Transport on the same corpora.

## 14. Reproducibility

The package contains source code, public corpora, generated corpora, benchmark scripts, result CSV/JSON files, figures, a citation checker, and tests [15]. The reviewer-facing smoke command is `python3 scripts/run_smoke_validation.py` and full validation is `python3 scripts/run_full_validation.py`. aioquic tests skip gracefully when aioquic is unavailable; installing `requirements-dev.txt` enables complete QUIC stream validation. Every table and figure in this paper is regenerated from the committed result files by `scripts/build_manuscript_v3_9.py` and `scripts/make_journal_figures_v2_8.py`. The framing-repricing and dictionary-delta baseline is regenerated by `benchmarks/run_framing_dictionary_baseline.py`, the miss-rate sensitivity sweep by `benchmarks/run_quic_miss_rate_sensitivity.py`, the QUIC confidence-interval table by `benchmarks/summarize_quic_statistical_evidence.py`, and the verifier-hardening properties of Section 4.5 (authentication-first ordering, bounded replay window, context binding under a valid tag) are pinned by dedicated unit tests. Each version of this manuscript is accompanied by an annotated git tag and a GitHub release whose assets include the final PDF and DOCX; the SHA-256 hashes of both files are pinned in `MANUSCRIPT_SHA256.txt` in the tagged tree, so reviewers can verify that the reviewed PDF matches the archived artifact.

## 15. Conclusion

ReduLink is best understood as authenticated, scoped reference substitution for encrypted endpoint streams. Its byte savings are conditional and are often reproduced or exceeded by simpler fixed-block reuse, compression, or rsync; the contribution is to make reference substitution explicit, individually authenticated, privacy-scoped, repairable, and compatible with native QUIC stream transport, and to position it precisely against both in-network redundancy elimination and HTTP shared-dictionary transport. The evidence supports ReduLink for warm-state object-aligned or layer-like transfers, while negative source-release results show where it should not be used.

## Data and Code Availability

All source code, public-corpus manifests, generated corpora, benchmark scripts, result CSV/JSON files, figures, the citation checker, and the test suite are openly available in the ReduLink repository [15] under the MIT License. Every figure and table is regenerated from the committed result files, so the reported numbers can be reproduced from the artifact. Large external public-release corpora are reconstructed from documented public sources rather than redistributed.

## References

- [1] N. T. Spring and D. Wetherall, 'A protocol-independent technique for eliminating redundant network traffic,' ACM SIGCOMM, 2000.
- [2] A. Anand, V. Sekar, and A. Akella, 'SmartRE: An architecture for coordinated network-wide redundancy elimination,' ACM SIGCOMM, 2009.
- [3] A. Anand et al., 'Redundancy in network traffic: Findings and implications,' ACM SIGMETRICS, 2009.
- [4] B. Aggarwal et al., 'EndRE: An end-system redundancy elimination service for enterprises,' USENIX NSDI, 2010.
- [5] A. Muthitacharoen, B. Chen, and D. Mazieres, 'A low-bandwidth network file system,' ACM SOSP, 2001.
- [6] J. Iyengar and M. Thomson, 'QUIC: A UDP-based multiplexed and secure transport,' RFC 9000, IETF, 2021.
- [7] M. Thomson and S. Turner, 'Using TLS to secure QUIC,' RFC 9001, IETF, 2021.
- [8] J. Iyengar and I. Swett, 'QUIC loss detection and congestion control,' RFC 9002, IETF, 2021.
- [9] M. Kuehlewind and B. Trammell, 'Applicability of the QUIC transport protocol,' RFC 9308, IETF, 2022.
- [10] B. Trammell et al., 'Manageability of the QUIC transport protocol,' RFC 9312, IETF, 2022.
- [11] W. Xia, X. Zou, Y. Zhou, H. Jiang, C. Liu, D. Feng, Y. Hua, Y. Hu, and Y. Zhang, 'The design of fast content-defined chunking for data deduplication based storage systems,' IEEE Transactions on Parallel and Distributed Systems, vol. 31, no. 9, 2020.
- [12] M. Gregoriadis, L. Balduf, B. Scheuermann, and J. Pouwelse, 'A thorough investigation of content-defined chunking algorithms for data deduplication,' arXiv:2409.06066, 2024.
- [13] B. Alexeev, C. Percival, and Y. X. Zhang, 'Chunking attacks on file backup services using content-defined chunking,' arXiv:2504.02095, 2025.
- [14] IEEE 802.3 Ethernet Working Group, 'IEEE 802.3 Ethernet Working Group active projects,' accessed June 2026.
- [15] M. Nguyen, 'ReduLink artifact and reproducibility package,' tagged release v3.9-journal-submission, GitHub, 2026.  
<https://github.com/pinkysworld/redulink-deduplex-quic/releases/tag/v3.9-journal-submission>
- [16] D. Harnik, B. Pinkas, and A. Shulman-Peleg, 'Side channels in cloud services: Deduplication in cloud storage,' IEEE Security and Privacy, 2010.
- [17] M. Bellare, S. Keelveedhi, and T. Ristenpart, 'DupLESS: Server-aided encryption for deduplicated storage,' USENIX Security, 2013.
- [18] aioquic project contributors, 'aioquic: QUIC and HTTP/3 implementation in Python,' software repository, version 1.3.0, 2026.  
<https://github.com/aiohttp/aioquic>
- [19] J. Mogul, B. Krishnamurthy, F. Douglass, A. Feldmann, Y. Goland, A. van Hoff, and D. Hellerstein, 'Delta encoding in HTTP,' RFC 3229, IETF, 2002.
- [20] D. Korn, J. MacDonald, J. Mogul, and K. Vo, 'The VCDIFF generic differencing and compression data format,' RFC 3284, IETF, 2002.
- [21] J. Butler, W.-H. Lee, B. McQuade, and K. Mixter, 'A proposal for shared dictionary compression over HTTP (SDCH),' IETF Internet-Draft, 2008.
- [22] P. Meenan and Y. Weiss, 'Compression dictionary transport,' RFC 9842, IETF, 2025.
- [23] H. Krawczyk, M. Bellare, and R. Canetti, 'HMAC: Keyed-hashing for message authentication,' RFC 2104, IETF, 1997.
- [24] H. Krawczyk and P. Eronen, 'HMAC-based extract-and-expand key derivation function (HKDF),' RFC 5869, IETF, 2010.
- [25] R. Jain, D.-M. Chiu, and W. Hawe, 'A quantitative measure of fairness and discrimination for shared computer systems,' DEC Research Report TR-301, 1984.
- [26] Y. Collet and M. Kucherawy, 'Zstandard compression and the application/zstd media type,' RFC 8878, IETF, 2021.